

```
In [1]: import numpy as np
import pandas as pd
from matplotlib.pyplot import subplots
```

```
In [3]: import statsmodels.api as sm
```

```
In [5]: from statsmodels.stats.outliers_influence import variance_inflation_factor as VIF
from statsmodels.stats.anova import anova_lm
```

```
In [7]: from ISLP import load_data
from ISLP.models import (
    ModelSpec as MS,
    summarize,
    poly
)
```

```
In [9]: dir()
```

```
Out[9]: ['In',
        'MS',
        'Out',
        'VIF',
        '_',
        '_2',
        '_4',
        '_6',
        '_8',
        '_',
        '_',
        '__builtin__',
        '__builtins__',
        '__doc__',
        '__loader__',
        '__name__',
        '__package__',
        'pandas',
        'session',
        'spec',
        'dh',
        'i',
        'i1',
        'i2',
        'i3',
        'i4',
        'i5',
        'i6',
        'i7',
        'i8',
        'i9',
        'ih',
        'ii',
        'iii',
        'oh',
        'anova_lm',
        'dataframe_columns',
        'dataframe_hash',
        'dtypes_str',
        'exit',
        'get_dataframes',
        'get_ipython',
        'getpass',
        'hashlib',
        'import_pandas_safely',
        'is_data_frame',
        'json',
        'load_data',
        'np',
        'open',
        'pd',
        'poly',
        'quit',
        'sm',
        'subplots',
        'summarize']
```

```
In [11]: A = np.array([3,5,11])
dir(A)
```

```
Out[11]: ['T',
          '__abs__',
          '__add__',
          '__and__',
          '__array__',
          '__array_finalize__',
          '__array_function__',
          '__array_interface__',
          '__array_prepare__',
          '__array_priority__',
          '__array_struct__',
          '__array_ufunc__',
          '__array_wrap__',
          '__bool__',
          '__buffer__',
          '__class__',
          '__class_getitem__',
          '__complex__',
          '__contains__',
          '__copy__',
          '__deepcopy__',
          '__delattr__',
          '__delitem__',
          '__dir__',
          '__divmod__',
          '__dlpack__',
          '__dlpack_device__',
          '__doc__',
          '__eq__',
          '__float__',
          '__floordiv__',
          '__format__',
          '__ge__',
          '__getattr__',
          '__getitem__',
          '__getstate__',
          '__gt__',
          '__hash__',
          '__iadd__',
          '__iand__',
          '__ifloordiv__',
          '__ilshift__',
          '__imatmul__',
          '__imod__',
          '__imul__',
          '__index__',
          '__init__',
          '__init_subclass__',
          '__int__',
          '__invert__',
          '__ior__',
          '__ipow__',
          '__irshift__',
          '__isub__',
          '__iter__',
          '__itruediv__',
          '__ixor__',
          '__le__',
          '__len__',
          '__lshift__',
```

```
'__lt__',
'__matmul__',
'__mod__',
'__mul__',
'__ne__',
'__neg__',
'__new__',
'__or__',
'__pos__',
'__pow__',
'__radd__',
'__rand__',
'__rdivmod__',
'__reduce__',
'__reduce_ex__',
'__repr__',
'__rfloordiv__',
'__rlshift__',
'__rmatmul__',
'__rmod__',
'__rmul__',
'__ror__',
'__rpow__',
'__rrshift__',
'__rshift__',
'__rsub__',
'__rtruediv__',
'__rxor__',
'__setattr__',
'__setitem__',
'__setstate__',
'__sizeof__',
'__str__',
'__sub__',
'__subclasshook__',
'__truediv__',
'__xor__',
'all',
'any',
'argmax',
'argmin',
'argpartition',
'argsort',
'astype',
'base',
'byteswap',
'choose',
'clip',
'compress',
'conj',
'conjugate',
'copy',
'ctypes',
'cumprod',
'cumsum',
'data',
'diagonal',
'dot',
'dtype',
'dump',
```

```
'dumps',
'fill',
'flags',
'flat',
'flatten',
'getfield',
'imag',
'item',
'itemset',
'itemsize',
'max',
'mean',
'min',
'nbytes',
'ndim',
'newbyteorder',
'nonzero',
'partition',
'prod',
'ptp',
'put',
'ravel',
'real',
'repeat',
'reshape',
'resize',
'round',
'searchsorted',
'setfield',
'setflags',
'shape',
'size',
'sort',
'squeeze',
'std',
'strides',
'sum',
'swapaxes',
'take',
'tobytes',
'tofile',
'tolist',
'tostring',
'trace',
'transpose',
'var',
'view']
```

```
In [13]: A.sum()
```

```
Out[13]: 19
```

```
In [15]: Boston = load_data("Boston")
Boston.columns
```

```
Out[15]: Index(['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad', 'tax',
               'ptratio', 'lstat', 'medv'],
              dtype='object')
```

```
In [17]: X = pd.DataFrame({
        'intercept': np.ones(Boston.shape[0]),
        'lstat': Boston['lstat']
    })

X[:4]
```

```
Out[17]:
```

	intercept	lstat
0	1.0	4.98
1	1.0	9.14
2	1.0	4.03
3	1.0	2.94

```
In [19]: y = Boston['medv']
        model = sm.OLS(y, X)
        results = model.fit()
```

```
In [21]: summarize(results)
```

```
Out[21]:
```

	coef	std err	t	P> t
intercept	34.5538	0.563	61.415	0.0
lstat	-0.9500	0.039	-24.528	0.0

```
In [23]: design = MS(['lstat'])
        design = design.fit(Boston)
        X = design.transform(Boston)
        X[:4]
```

```
Out[23]:
```

	intercept	lstat
0	1.0	4.98
1	1.0	9.14
2	1.0	4.03
3	1.0	2.94

```
In [25]: #fit_transform() method#
        design = MS(['lstat'])
        X = design.fit_transform(Boston)

X[:4]
```

```
Out[25]:
```

	intercept	lstat
0	1.0	4.98
1	1.0	9.14
2	1.0	4.03
3	1.0	2.94

```
In [27]: results.summary()
```

Out[27]: OLS Regression Results

Dep. Variable:	medv	R-squared:	0.544			
Model:	OLS	Adj. R-squared:	0.543			
Method:	Least Squares	F-statistic:	601.6			
Date:	Mon, 17 Aug 2026	Prob (F-statistic):	5.08e-88			
Time:	09:24:02	Log-Likelihood:	-1641.5			
No. Observations:	506	AIC:	3287.			
Df Residuals:	504	BIC:	3295.			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t 	[0.025	0.975]
intercept	34.5538	0.563	61.415	0.000	33.448	35.659
lstat	-0.9500	0.039	-24.528	0.000	-1.026	-0.874
Omnibus:	137.043	Durbin-Watson:	0.892			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	291.373			
Skew:	1.453	Prob(JB):	5.36e-64			
Kurtosis:	5.319	Cond. No.	29.7			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly speci

```
In [29]: results.params
```

```
Out[29]: intercept    34.553841
lstat              -0.950049
dtype: float64
```

```
In [31]: #The get_prediction()#
new_df = pd.DataFrame({'lstat':[5, 10, 15]})
newX = design.transform(new_df)
newX
```

```
Out[31]:   intercept  lstat
0         1.0      5
1         1.0     10
2         1.0     15
```

```
In [33]: new_predictions = results.get_prediction(newX);
new_predictions.predicted_mean
```

```
Out[33]: array([29.80359411, 25.05334734, 20.30310057])
```

```
In [35]: new_predictions.conf_int(alpha=0.05)
```

```
Out[35]: array([[29.00741194, 30.59977628],  
               [24.47413202, 25.63256267],  
               [19.73158815, 20.87461299]])
```


```
In [37]: #Prediction intervals#  
new_predictions.conf_int(obs=True , alpha=0.05)
```

```
Out[37]: array([[17.56567478, 42.04151344],  
               [12.82762635, 37.27906833],  
               [ 8.0777421 , 32.52845905]])
```

```
In [39]: #Defining Functions#  
def abline(ax, b, m):  
    """Add a line with slope m and intercept b to ax."""  
    xlim = ax.get_xlim()  
    ylim = [m * xlim[0] + b, m * xlim[1] + b]  
    ax.plot(xlim, ylim)
```

```
In [41]: def abline(ax, b, m, *args, **kwargs):  
    """Add a line with slope m and intercept b to ax."""  
    xlim = ax.get_xlim()  
    ylim = [m * xlim[0] + b, m * xlim[1] + b]  
    ax.plot(xlim, ylim, *args, **kwargs)
```


```
In [49]: ax = Boston.plot.scatter('lstat', 'medv')  
  
abline(  
    ax,  
    results.params.iloc[0],  
    results.params.iloc[1],  
    'r--',  
    linewidth=3  
)
```

In [53]:


```
#Linear Regression#  
ax = subplots(figsize=(8, 8))[1]  
  
ax.scatter(results.fittedvalues, results.resid)  
ax.set_xlabel('Fitted value')  
ax.set_ylabel('Residual')  
ax.axhline(0, c='k', ls='--')
```

Out[53]: <matplotlib.lines.Line2D at 0x1cbaaed1340>



```
In [55]: infl = results.get_influence()
ax = subplots(figsize=(8, 8))[1]
ax.scatter(
    np.arange(X.shape[0]),
    infl.hat_matrix_diag
)
ax.set_xlabel('Index')
ax.set_ylabel('Leverage')
np.argmax(infl.hat_matrix_diag)
```

Out[55]: 374



```
In [57]: #Multiple Linear Regression#
X = MS(['lstat', 'age']).fit_transform(Boston)
model1 = sm.OLS(y, X)
results1 = model1.fit()
summarize(results1)
```

```
Out[57]:
```

	coef	std err	t	P> t
intercept	33.2228	0.731	45.458	0.000
lstat	-1.0321	0.048	-21.416	0.000
age	0.0345	0.012	2.826	0.005

```
In [63]: terms = Boston.columns.drop('medv')
```

```
In [65]: Boston.columns
```

```
Out[65]: Index(['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad', 'tax',
                'ptratio', 'lstat', 'medv'],
                dtype='object')
```

```
In [67]: X = MS(terms).fit_transform(Boston)
model = sm.OLS(y, X)
results = model.fit()
summarize(results)
```

Out[67]:

	coef	std err	t	P> t
intercept	41.6173	4.936	8.431	0.000
crim	-0.1214	0.033	-3.678	0.000
zn	0.0470	0.014	3.384	0.001
indus	0.0135	0.062	0.217	0.829
chas	2.8400	0.870	3.264	0.001
nox	-18.7580	3.851	-4.870	0.000
rm	3.6581	0.420	8.705	0.000
age	0.0036	0.013	0.271	0.787
dis	-1.4908	0.202	-7.394	0.000
rad	0.2894	0.067	4.325	0.000
tax	-0.0127	0.004	-3.337	0.001
ptratio	-0.9375	0.132	-7.091	0.000
lstat	-0.5520	0.051	-10.897	0.000

```
In [69]: minus_age = Boston.columns.drop(['medv', 'age'])
Xma = MS(minus_age).fit_transform(Boston)
model1 = sm.OLS(y, Xma)
summarize(model1.fit())
```

Out[69]:

	coef	std err	t	P> t
intercept	41.5251	4.920	8.441	0.000
crim	-0.1214	0.033	-3.683	0.000
zn	0.0465	0.014	3.379	0.001
indus	0.0135	0.062	0.217	0.829
chas	2.8528	0.868	3.287	0.001
nox	-18.4851	3.714	-4.978	0.000
rm	3.6811	0.411	8.951	0.000
dis	-1.5068	0.193	-7.825	0.000
rad	0.2879	0.067	4.322	0.000
tax	-0.0127	0.004	-3.333	0.001
ptratio	-0.9346	0.132	-7.099	0.000
lstat	-0.5474	0.048	-11.483	0.000

```
In [71]: #Multivariate Goodness of Fit#
#List Comprehension#
vals = [VIF(X, i)
        for i in range(1, X.shape[1])]

vif = pd.DataFrame(
    {'vif': vals},
```

```

        index=X.columns[1:]
    )

vif

```

Out[71]:

	vif
crim	1.767486
zn	2.298459
indus	3.987181
chas	1.071168
nox	4.369093
rm	1.912532
age	3.088232
dis	3.954037
rad	7.445301
tax	9.002158
ptratio	1.797060
lstat	2.870777

```

In [73]: vals = []

for i in range(1, X.values.shape[1]):
    vals.append(VIF(X.values, i))

```

```

In [75]: #Interaction Terms#
X = MS([
    'lstat',
    'age',
    ('lstat', 'age')
]).fit_transform(Boston)
model2 = sm.OLS(y, X)
summarize(model2.fit())

```

Out[75]:

	coef	std err	t	P> t
intercept	36.0885	1.470	24.553	0.000
lstat	-1.3921	0.167	-8.313	0.000
age	-0.0007	0.020	-0.036	0.971
lstat:age	0.0042	0.002	2.244	0.025

```

In [77]: #Non-linear Transformations of the Predictors#
X = MS([
    poly('lstat', degree=2),
    'age'
]).fit_transform(Boston)
model3 = sm.OLS(y, X)
results3 = model3.fit()
summarize(results3)

```


	coef	std err	t	P> t
intercept	17.7151	0.781	22.681	0.0
poly(lstat, degree=2)[0]	-179.2279	6.733	-26.620	0.0
poly(lstat, degree=2)[1]	72.9908	5.482	13.315	0.0
age	0.0703	0.011	6.471	0.0

```
In [79]: anova_lm(results1 , results3)
```

	df_resid	ssr	df_diff	ss_diff	F	Pr(>F)
0	503.0	19168.128609	0.0	NaN	NaN	NaN
1	502.0	14165.613251	1.0	5002.515357	177.278785	7.468491e-35

```
In [81]: #Linear Regression#
ax = subplots(figsize=(8, 8))[1]
ax.scatter(results3.fittedvalues, results3.resid)
ax.set_xlabel('Fitted value')
ax.set_ylabel('Residual')
ax.axhline(0, c='k', ls='--')
```

```
Out[81]: <matplotlib.lines.Line2D at 0x1cbab0b8260>
```



```
In [83]: #Qualitative Predictors#
Carseats = load_data('Carseats')
Carseats.columns
```

```
Out[83]: Index(['Sales', 'CompPrice', 'Income', 'Advertising', 'Population', 'Price',
               'ShelveLoc', 'Age', 'Education', 'Urban', 'US'],
              dtype='object')
```

```
In [85]: allvars = list(Carseats.columns.drop('Sales'))

y = Carseats['Sales']
final = allvars + [
    ('Income', 'Advertising'),
    ('Price', 'Age')
]
X = MS(final).fit_transform(Carseats)
model = sm.OLS(y, X)
summarize(model.fit())
```

```
Out[85]:
```

	coef	std err	t	P> t
intercept	6.5756	1.009	6.519	0.000
CompPrice	0.0929	0.004	22.567	0.000
Income	0.0109	0.003	4.183	0.000
Advertising	0.0702	0.023	3.107	0.002
Population	0.0002	0.000	0.433	0.665
Price	-0.1008	0.007	-13.549	0.000
ShelveLoc[Good]	4.8487	0.153	31.724	0.000
ShelveLoc[Medium]	1.9533	0.126	15.531	0.000
Age	-0.0579	0.016	-3.633	0.000
Education	-0.0209	0.020	-1.063	0.288
Urban[Yes]	0.1402	0.112	1.247	0.213
US[Yes]	-0.1576	0.149	-1.058	0.291
Income:Advertising	0.0008	0.000	2.698	0.007
Price:Age	0.0001	0.000	0.801	0.424

```
In [ ]:
```